POLITECNICO
MILANO 1863

Artificial Neural Networks and Deep Learning
Homework 1 - Image Classification

Frantuma Elia - 10567359 - 945729,
Fucci Tiziano - 10524029 - 946638

A.Y. 2020/2021

Table of contents

Chapter 1

Introduction

1.1 Description of the task

The homework consists in an image classification problem on the proposed dataset. In particular, it is required to classify images depicting groups of people based on the number of masked people. In the specific, the solution must discriminate between images depending on the following cases:

1. no person in the image is wearing a mask;

2. all the people in the image are wearing a mask;

3. someone in the image is not wearing a mask.

In the following figure, one sample image for each class is shown.

(1)   (2)   (3)

Thus, the classification is performed on 3 different classes. Being a classification problem, given an image, the goal is to predict the correct class label.

1.2 Dataset

The dataset is composed by .jpg images of different size, divided in two folders:

- training (5614 images)

- test (450 images)

A JSON file, containing the labels of the training images, is attached to the dataset. The dataset requires the training images to be divided into three folders (created manually), corresponding to the three target classes. The division was done with a python script, which is contained in the notebooks.

Note: the dataset path expected by the script is:
..\\artificial-neural-networks-and-deep-learning-2020\\MaskDataset

1.2.1 Data augmentation

We have performed data augmentation in order to increase the dataset dimension. Some of the parameters used to perform the transformations are: brightness, zoom, horizontal/vertical shift and flip. We didn't use rotation, since we noticed that the filling could reduce the performance of the classifier.

1.3 Validation set

No automatic validation set is provided. This means that a subset of the training set must be used to perform validation.

In our case, we parametrized the number of training images to be moved into the validation set, with values between 3-15%.

1.4 Test set

The test set was left untouched, apart from the creation of one directory, named "unlabeled", which contains all the images.

1.5 Evaluation

Submissions are evaluated on multiclass accuracy, which is simply the average number of observations with the correct label. To submit a prediction, it is necessary to produce a .csv file containing the predictions associated to

each test image. This file has to be submitted to the Kaggle page of the competition to obtain the accuracy score on the test set.

Chapter 2

Neural network architecture

2.1 Ex-novo architecture

First, we tried to build a sequential neural network from scratch: we made many experiments, each one with a slightly different parameter configuration. Among them:

- `start_f`: the number of initial filters;

- `depth`: the number of convolutional layers;

- the learning rate;

- the batch size;

- the size of the input image;

- the shape of the neuron activation function;

- the number of fully connected layers.

After few experiments, we found that the best model was a convolutional neural network having `start_f = 9` and `depth = 7`. The complete model of the network is described in the attached file `HomeworkImages.ipynb`.

2.1.1 Score

The best score of the network was 0.85555, obtained with the following settings:

- `start_f = 9`;

- `depth = 7`;

- image resolution: 348x522;

- batch size = 8.

2.1.2   Diagrams

Here we show how loss and accuracy changed during the network training, both for the training (grey) and the validation set (orange).


Figure 2.1: Accuracy plot

Figure 2.2: Loss plot

2.2   Transfer learning with VGG

The final version of the ex-novo architecture was pretty accurate, but we could not achieve better results by just changing the parameters. We decided to use a pre-trained, reliable and successful architecture.
Our first try was with VGG16, because we've seen it during the lectures. Also in this case we made many experiments, each one with a different parameter configuration:

- `freeze_until`: the layer of the network from which we want to fine-tune;

- the learning rate;

- the size of the input image;

- the shape of the neuron activation function;

- the number of fully connected layers.

- the number of neurons in fully connected layers.

2.2.1   Score

The best score of the network was 0.89333, obtained with the following settings:

- `freeze_until = 11;`

- one fully connnected layer with 32 neurons before the softmax layer.

2.2.2   Diagrams

Here we show how loss and accuracy changed during the network training, both for the training (grey) and the validation set (orange).

Figure 2.3: Accuracy plot

Figure 2.4: Loss plot

2.3 Transfer learning with ImageResNetV2

For our last attempt we decided to use ImageResNetV2, the best performing on the ImageNet dataset. The approach in this case has been similar to the one with VGG, since we only changed the backbone. One significant change that has been made is the replacement of the fully connected layer(s) with a Global Average Pooling layer.

2.3.1 Score

The best score of the network was 0.96888, obtained with the following settings:

- `freeze_until = 150;`

- image resolution: 348x522;

- batch size = 8.

2.3.2 Plots

Here we show how loss and accuracy changed during the network training, both for the training (grey) and the validation set (orange).


Figure 2.5: Accuracy plot

Figure 2.6: Loss plot

Chapter 3

References

3.1   Links

- GitHub repository of the project: `https://github.com/tizianofucci/A2NDLKaggle`